

1. Source code

```
# stock_price_prediction/app.py

import gradio as gr
import pandas as pd
import numpy as np
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import LSTM, Dense
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
import yfinance as yf
import os

# Load or train model
MODEL_PATH = "stock_model.h5"
scaler = MinMaxScaler()

def get_stock_data(ticker="AAPL"):
    df = yf.download(ticker, period="2y")
```

```
df = df[['Close']].dropna() return
df
```

```
def create_dataset(data, time_step=60):
    X, y = [], []
    for i in range(time_step, len(data)):
        X.append(data[i-time_step:i, 0])
        y.append(data[i, 0])
    return np.array(X), np.array(y)
```

```
def train_model(df):
    data = scaler.fit_transform(df.values)
    X, y = create_dataset(data)
    X = X.reshape(X.shape[0], X.shape[1], 1)

    model = Sequential()
    model.add(LSTM(units=50, return_sequences=True, input_shape=(X.shape[1],
1)))
    model.add(LSTM(units=50))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mean_squared_error')

    model.fit(X, y, epochs=5, batch_size=32, verbose=1)
    model.save(MODEL_PATH)
    return model
```

```
def load_or_train_model(df):
    if os.path.exists(MODEL_PATH):
        return load_model(MODEL_PATH)
    else:
        return train_model(df)
```

```
def prepare_input_sequence(df, window=60):
    data = scaler.transform(df.values)
    return np.reshape(data[-window:], (1, window, 1))
```

```
def predict_prices(days, ticker):
    df = get_stock_data(ticker)
    model = load_or_train_model(df)
    input_seq = prepare_input_sequence(df)

    predictions = []
    for _ in range(days):
        pred = model.predict(input_seq, verbose=0)[0, 0]
```

```

        predictions.append(pred)
        input_seq = np.append(input_seq[0][1:],
                               [[pred]], axis=0) input_seq =
        input_seq.reshape(1, input_seq.shape[0], 1)

    return scaler.inverse_transform(np.array(predictions).reshape(-1,

1)).flatten()
def predict_and_plot(days, ticker):
    preds = predict_prices(days, ticker)
    days_range = list(range(1, days + 1))

    plt.figure(figsize=(8, 4))
    plt.plot(days_range, preds, marker='o', linestyle='-',
             color='blue')
    plt.title(f'Predicted Stock Prices for {ticker.upper()}')
    plt.xlabel("Days Ahead")
    plt.ylabel("Price (USD)")
    plt.grid(True)

    return plt.gcf(), {f'Day {i+1}': f'${p:.2f}' for i, p in
enumerate(preds)} # Gradio Interface
demo = gr.Interface(
    fn=predict_and_plot,
    inputs=[
        gr.Slider(1, 10, step=1, label="Days Ahead"),
        gr.Textbox(value="AAPL", label="Stock Ticker")
    ],
    outputs=["plot", "label"],
    title="Stock Price Prediction App",
    description="Predict the next few days of stock closing prices using an LSTM model trained on historical data."
)

demo.launch()

```